

Experiment No. 3
Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model
Date of Performance:
Date of Submission:

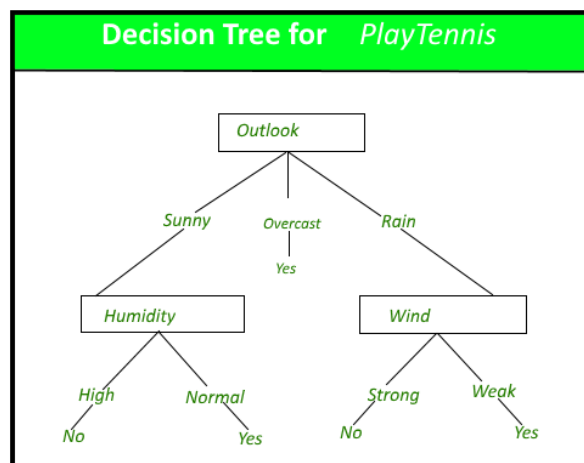


Aim: Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model.

Objective: To perform various feature engineering tasks, apply Decision Tree Algorithm on the given dataset and maximize the accuracy, Precision, Recall, F1 score. Improve the performance by performing different data engineering and feature engineering tasks.

Theory:

Decision Tree is the most powerful and popular tool for classification and prediction. A Decision tree is a flowchart-like tree structure, where each internal node denotes a test on an attribute, each branch represents an outcome of the test, and each leaf node (terminal node) holds a class label.



Dataset:

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:

>50K, <=50K.



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.

fnlwgt: continuous.

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.

education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

sex: Female, Male.

capital-gain: continuous.

capital-loss: continuous.

hours-per-week: continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala,



Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands.

Code:

Conclusion:

1. Discuss about the how categorical attributes have been dealt with during data pre-processing.1.

Categorical attributes in the dataset are workclass, education, marital status, relationship, race, sex, native country. These attributes are handled using the label encoder imported from sklearn.preprocessing import LabelEncoder

2. Discuss the hyper-parameter tuning done based on the decision tree obtained.

Hyper parameter tuning is not performed in this experiment but Hyperparameter tuning for decision trees involves optimizing the settings or parameters that are not learned from the data but are set prior to the training process. These hyperparameters can significantly impact the performance of decision trees.

3. Comment on the accuracy, confusion matrix, precision, recall and F1 score obtained.

The model's accuracy is decent at 81.18%, indicating a good overall classification performance. However, the precision score of 59.75% suggests that there may be a significant number of false positives. The recall of 62.00% is relatively balanced, indicating the model's ability to capture actual positives. The F1 score, which combines precision and recall, is 60.86%, showing a reasonable trade-off between precision and recall. Further analysis is needed to understand the specific trade-offs and context of these metrics.



Vidyavardhini's College of Engineering & Technology
Department of Computer Engineering

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import linear_model
from sklearn.tree import DecisionTreeClassifier
df= pd.read_csv('adult.csv')

df= pd.read_csv('adult.csv')
df.head()
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	rela
0	90	?	77053	HS-grad	9	Widowed	?	Nc
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Nc
2	66	?	186061	Some-college	10	Widowed	?	
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	

```
df.shape

(32561, 15)

df.describe()
```

	age	fnlwgt	education.num	capital.gain	capital.loss	hours.p
count	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000	32561
mean	38.581647	1.897784e+05	10.080679	1077.648844	87.303830	40
std	13.640433	1.055500e+05	2.572720	7385.292085	402.960219	12
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000	40
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	40
75%	48.000000	2.370510e+05	12.000000	0.000000	0.000000	45
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	99

```
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                    32561 non-null  int64
1   workclass              32561 non-null  object
2   fnlwgt                 32561 non-null  int64
3   education              32561 non-null  object
4   education.num          32561 non-null  int64
5   marital.status         32561 non-null  object
6   occupation              32561 non-null  object
7   relationship           32561 non-null  object
8   race                   32561 non-null  object
9   sex                    32561 non-null  object
10  capital.gain            32561 non-null  int64
11  capital.loss            32561 non-null  int64
12  hours.per.week          32561 non-null  int64
13  native.country          32561 non-null  object
14  income                  32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB

df['income'].value_counts()

<=50K    24720
>50K      7841
Name: income, dtype: int64
```

```
df['sex'].value_counts()
```

```
Male      21790
Female    10771
Name: sex, dtype: int64
```

```
df['native.country'].value_counts()
```

```
United-States    29170
Mexico           643
?                583
Philippines      198
Germany          137
Canada           121
Puerto-Rico     114
El-Salvador      106
India            100
Cuba             95
England          90
Jamaica          81
South            80
China            75
Italy            73
Dominican-Republic 70
Vietnam          67
Guatemala        64
Japan            62
Poland           60
Columbia         59
Taiwan           51
Haiti            44
Iran             43
Portugal         37
Nicaragua        34
Peru             31
Greece           29
France           29
Ecuador          28
Ireland          24
Hong             20
Cambodia         19
Trinidad&Tobago  19
Laos             18
Thailand         18
Yugoslavia       16
Outlying-US(Guam-USVI-etc) 14
Hungary          13
Honduras         13
Scotland         12
Holand-Netherlands 1
Name: native.country, dtype: int64
```

```
df['workclass'].value_counts()
```

```
Private      22696
Self-emp-not-inc 2541
Local-gov    2093
?            1836
State-gov    1298
Self-emp-inc 1116
Federal-gov  960
Without-pay  14
Never-worked 7
Name: workclass, dtype: int64
```

```
df['occupation'].value_counts()
```

```
Prof-specialty      4140
Craft-repair        4099
Exec-managerial     4066
Adm-clerical        3770
Sales               3650
Other-service       3295
Machine-op-inspct   2002
?                  1843
Transport-moving    1597
Handlers-cleaners   1370
Farming-fishing     994
Tech-support        928
Protective-serv     649
Priv-house-serv     149
Armed-Forces        9
Name: occupation, dtype: int64
```

```
df.replace('?', np.NaN, inplace = True)
df.head()
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship
0	90	NaN	77053	HS-grad	9	Widowed	NaN	Never-married
1	82	Private	132870	HS-grad	9	Widowed	Executive-managerial	Never-married
2	66	NaN	186061	Some-college	10	Widowed	NaN	Never-married
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Married
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Married

```
#fills missing values (NaNs) with the most recent non-missing value
df.fillna(method = 'ffill', inplace = True)
```

```
from sklearn.preprocessing import LabelEncoder
```

```
le = LabelEncoder()
df['workclass'] = le.fit_transform(df['workclass'])
df['education'] = le.fit_transform(df['education'])
df['marital.status'] = le.fit_transform(df['marital.status'])
df['occupation'] = le.fit_transform(df['occupation'])
df['relationship'] = le.fit_transform(df['relationship'])
df['race'] = le.fit_transform(df['race'])
df['sex'] = le.fit_transform(df['sex'])
df['native.country'] = le.fit_transform(df['native.country'])
df['income'] = le.fit_transform(df['income'])
```

```
df.head()
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship
0	90	8	77053	11	9	6	14	Never-married
1	82	3	132870	11	9	6	3	Never-married
2	66	3	186061	15	10	6	3	Never-married
3	54	3	140359	5	4	0	6	Married
4	41	3	264663	15	10	5	9	Married

```
df.describe()
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship
count	32561.000000	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000	32561.000000	32561.000000
mean	38.581647	3.102761	1.897784e+05	10.298210	10.080679	2.572720	2.572720	2.572720
std	13.640433	1.136995	1.055500e+05	3.870264	2.572720	1.000000	1.000000	1.000000
min	17.000000	0.000000	1.228500e+04	0.000000	1.000000	0.000000	0.000000	0.000000
25%	28.000000	3.000000	1.178270e+05	9.000000	9.000000	2.000000	2.000000	2.000000
50%	37.000000	3.000000	1.783560e+05	11.000000	10.000000	2.000000	2.000000	2.000000
75%	48.000000	3.000000	2.370510e+05	12.000000	12.000000	4.000000	4.000000	4.000000
max	90.000000	8.000000	1.484705e+06	15.000000	16.000000	6.000000	6.000000	6.000000

```
#Splitting the data set into features and outcome
X = df.drop(['income'], axis=1)
Y = df['income']
```

```
df.isnull().sum()
```

```
age      0
workclass 0
fnlwgt   0
education 0
education.num 0
marital.status 0
```



```

occupation      0
relationship    0
race            0
sex            0
capital.gain    0
capital.loss    0
hours.per.week  0
native.country  0
income          0
dtype: int64

```

```
df.duplicated().sum()
```

```
24
```

```
df=df.dropna()
```

```
#Splitting the data into test data and training data
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2, random_state = 42)
```

```
X_train.head()
```

	age	workclass	fnlwgt	education	education.num	marital.status	occupation
5514	26	3	256263	11	9	4	2
19777	24	3	170277	11	9	4	7
10781	36	3	75826	9	13	0	0
32240	22	6	24395	15	10	2	0
9876	31	1	356689	9	13	2	9

```
Y_train = Y_train.replace((np.inf, -np.inf, np.nan), 0).reset_index(drop=True)
```

```
Y_train.head()
```

```

0    0
1    0
2    0
3    0
4    0
Name: income, dtype: int64

```

```
Y_test = Y_test.replace((np.inf, -np.inf, np.nan), 0).reset_index(drop=True)
```

```

from sklearn.tree import DecisionTreeClassifier
dec_tree = DecisionTreeClassifier(random_state=42)
dec_tree.fit(X_train, Y_train)
Y_pred_dec_tree = dec_tree.predict(X_test)

```

```

import matplotlib.pyplot as plt
from sklearn.tree import DecisionTreeClassifier, plot_tree
plot_tree(dec_tree)
plt.show()

```



```
from sklearn.metrics import accuracy_score
from sklearn.metrics import f1_score
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score
```



```
print('Decision Tree Classifier:')
print('Accuracy score:',accuracy_score(Y_test, Y_pred_dec_tree) * 100)
print('Precision :',precision_score(Y_test,Y_pred_dec_tree) *100)
print('Recall: ',recall_score(Y_test,Y_pred_dec_tree)* 100)
print('F1 score: ',f1_score(Y_test,Y_pred_dec_tree) *100)
```

```
Decision Tree Classifier:
Accuracy score: 81.1761093198219
Precision : 59.749216300940446
Recall: 62.0039037085231
F1 score: 60.85568326947638
```